

# Lecture 05: Language modelling

# OVERVIEW

- What is language modeling
- Why do we need language modeling
- Probabilistic Language Models
- Language modeling with n-grams

# LANGUAGE MODEL

## Formal definition

- Given a finite vocabulary  $V$  of words (tokens):
- Formal language: let  $\Omega$  be a set of sequences of words from  $V$ ,  
 $\forall x \in \Omega, x$  is called a sentence
- Language Model: A probability distribution over the formal language  $\Omega$ ,

$$P : \Omega \rightarrow [0, 1] \quad \sum_{x \in \Omega} P(x) = 1$$

- A **Language Modeling** is a probability distribution over a sequence of words (tokens).

# LANGUAGE MODEL

A language model aims to answer the question which sequences are more likely?

*I would like to eat.*

*I would like eat to.*

*I like would eat to.*

*I to like would eat.*

We expect that regular and grammatically correct sentences will occur more often in text and speech than other weird sequences

# WHY DO WE NEED LANGUAGE MODELS?

Many NLP tasks require natural language output estimated from a sentence probability

- Machine translation: *"vents violents ce soir"*

$$P(\textit{high winds tonight}) > P(\textit{large winds tonight})$$

- Speech recognition: return a transcript of what was spoken

$$P(\textit{I saw a van}) \gg P(\textit{eyes awe of an})$$

- Spell-correction: *"the office is about fifteen minuets from my house"*

$$P(\textit{about fifteen minutes from}) > P(\textit{about fifteen minuets from})$$

- Summarization, question answering, etc

# PROBABILISTIC LANGUAGE MODEL

- Probabilistic language models are based on the grouping of words into chunks called n-grams in order to estimate

$$P(\mathbf{x}^{(t+1)} \mid \mathbf{x}^{(t)}, \dots, \mathbf{x}^{(1)})$$

- The main idea:
  - collect statistics about the frequency of different *n-grams*
    - n-gram probabilities are estimated by counting the frequency of occurrence in the vocabulary
  - use this to estimate the next word given a history of observed words

# LANGUAGE MODELING

*How do we compute the probability of a sequence of tokens  $x^1, x^2, \dots, x^{t-1}, x^t$  from our vocabulary  $V$ ?*

The probability of a sentence  $x^1, x^2, \dots, x^t$  is the joint probability  $P(x^1, x^2, \dots, x^t)$  and estimate by the chain rule of probability

$$\begin{aligned} P(x^1, x^2, \dots, x^{t-1}, x^t) &= P(x^1 | x^2, \dots, x^{t-1}, x^t) P(x^2, \dots, x^{t-1}, x^t) \\ &= P(x^1) P(x^2 | x^1) P(x^3 | x^1, x^2) \dots P(x^t | x^1, \dots, x^{t-1}) \\ &= P(x^1) \prod_{k=1}^t P(x^k | x^1, \dots, x^{k-1}) \end{aligned}$$

# LANGUAGE MODEL

Given the sentence “*the student open their books*”

How do we compute the probability of  $P(\textit{the student open their books})$

$$\begin{aligned} P(\textit{the student open their books}) = & P(\textit{books}|\textit{the student open their}) * \\ & P(\textit{their}|\textit{the student open}) * \\ & P(\textit{open}|\textit{the student}) * \\ & P(\textit{student}|\textit{the}) * \\ & P(\textit{the}) \end{aligned}$$

We estimate the probabilities by counting

$$P(\textit{books}|\textit{the student open their}) = \frac{\text{count}(\textit{the student open their books})}{\text{count}(\textit{the student open their})}$$



# LANGUAGE MODEL – ESTIMATING PROBABILITIES

A language model is a distribution over all word-sequences  $x_1x_2, \dots, x_n$  in a vocabulary  $V$

$$\sum_{\langle x_1x_2, \dots, x_n \rangle} P(x_1x_2, \dots, x_n) = 1$$

by derivation is given by  $P(x^1, x^2, \dots, x^{t-1}, x^t) = P(x^1) \prod_{k=1}^t P(x^k | x^1, \dots, x^{k-1})$

To estimate the probability of each sequence we need to

- first estimate  $P(x^1)$
- Estimate probabilities  $P(x^k | x^1, \dots, x^{k-1})$  for all  $x^1, \dots, x^k$

# LANGUAGE MODEL – ESTIMATING PROBABILITIES

- Relative frequency from a corpus

$$\begin{aligned} P(x_i | x_1, \dots, x_{i-1}) &= \frac{\text{count}(x_1, \dots, x_{i-1}, x_i)}{\text{count}(x_1, \dots, x_{i-1})} \\ &= \frac{\text{count}(x_1, \dots, x_{i-1}, x_i)}{\sum_{x \in V} \text{count}(x_1, \dots, x_{i-1}, x)} \end{aligned}$$


NB: this is the estimate for all sequences of length /

- Suppose  $|V| = 1000$ , all sentences are approximately 10 word long, then we need to estimate  $1000^{10}$  probabilities.
- no corpus is large enough to obtain an unbiased estimate of the probabilities

# LANGUAGE MODELING

**Language Modeling** can be reduced to the task of predicting the next

$$P(x^1, x^2, \dots, x^{t-1}, x^t) = P(x^1) \prod_{n=2}^t P(x^n | x^1, \dots, x^{n-1})$$



Given a sequence of words  $x^{(1)}, x^{(2)}, \dots, x^{(t)}$  compute the probability distribution of the next word  $x^{(t+1)}$ :

$$P(\mathbf{x}^{(t+1)} | \mathbf{x}^{(t)}, \dots, \mathbf{x}^{(1)})$$

# LANGUAGE MODELING – MARKOV ASSUMPTION

Markov assumption

- Independent and identical trials
- There is a fixed and finite  $k$  such that all word depends only on the preceding  $k-1$  words

$$P(x_{i+1}|x_1, \dots, x_i) \approx P(x_{i+1}|x_{i-k}, \dots, x_i) \quad \forall k \geq 0$$

- Model: an  $k^{th}$ - order Markov model
- n-gram: statistics of an k-order Markov model is  $k + 1 - gram$  model

$$P(x_i | x_{i-k}, \dots, x_{i-1}) = \frac{\text{count}(x_{i-k}, \dots, x_{i-1}, x_i)}{\sum_{x \in V} \text{count}(x_{i-k}, \dots, x_{i-1}, x)}$$

# LANGUAGE MODEL – MARKOV ASSUMPTION AND N-GRAMS

- N-gram language models are based on probabilities of chunks of word.

*The student open their*

- Definition: An n-gram is a chunk of n consecutive words.
  - *unigram*: unit of single word – “the”, “student”, “opened”, “their”
  - *bigrams*: unit of double words – “the student”, “student opened”, “opened their”
  - *trigram*: unit of triple words – “the student opened”, “student opened their”
  - *4-gram*: unit of 4 words – “the student opened their”
- The main idea behind *n-gram* models is to collect statistics about the frequency of different *n-grams* and use this to predict the next word.

# LANGUAGE MODELING – MARKOV ASSUMPTION

The order of a Markov model is defined by the length of its history or *n-gram* ( $n = k+1$ )

$0^{th}$  – order:  $P(x_1, \dots, x_n) \approx P(x_1) \prod_{i=1}^{n-1} P(x_{i+1}) \approx P(x_1)P(x_2) \dots P(x_n)$   
|history| = 0

$1^{st}$  – order:  $P(x_1, \dots, x_n) \approx P(x_1) \prod_{i=1}^{n-1} P(x_{i+1}|x_i)$   
|history| = 1

$2^{nd}$  – order:  $P(x_1, \dots, x_n) \approx P(x_1) \prod_{i=1}^{n-1} P(x_{i+1}|x_i, x_{i-1})$   
|history| = 2

$k^{th}$  – order:  $P(x_1, \dots, x_n) \approx P(x_1) \prod_{i=1}^{n-1} P(x_{i+1}|x_i, \dots, x_{i-k})$   
|history| = k

# FROM N-GRAM PROBABILITIES TO LANGUAGE MODEL

In a trigram model

$$P(x^1x^2x^3) = P(x^1)P(x^2|x^1)P(x^3|x^1x^2)$$

- The only trigram is  $P(x^3|x^1x^2)$ 
  - $P(x^1)$  and  $P(x^2|x^1)$  are not trigrams thus are from a different probability distribution
- **Solution:**
  - add  $n-1$  beginning of sentence ( $\langle s \rangle$ ) symbols  
 $\langle s \rangle \langle s \rangle x^1 x^2 x^3 \dots$
  - similarly add  $n-1$  end of sentence symbols  
 $\dots x^1 x^2 x^3 \langle /s \rangle \langle /s \rangle$

# ESTIMATING BIGRAM PROBABILITIES

Bigram estimate of the probability of the sentence *“I want English food”*

$$\begin{aligned} P(<s> \text{ I want english food } </s>) &= \\ P(\text{I} | <s>) & \\ \times P(\text{want} | \text{I}) & \\ \times P(\text{english} | \text{want}) & \\ \times P(\text{food} | \text{english}) & \\ \times P(</s> | \text{food}) & \\ = .000031 & \end{aligned}$$

How is it estimated

- $P(\text{english} | \text{want}) = .0011$
- $P(\text{chinese} | \text{want}) = .0065$
- $P(\text{to} | \text{want}) = .66$
- $P(\text{eat} | \text{to}) = .28$
- $P(\text{food} | \text{to}) = 0$
- $P(\text{want} | \text{spend}) = 0$
- $P(\text{i} | <s>) = .25$



# FROM N-GRAM PROBABILITIES TO LANGUAGE MODEL

- With the start or end of sentence token(s) we define a new vocabulary

$$V^* = V \cup \{</s>\}$$

or

$$V^* = V \cup \{< s >\}$$

- With the new vocabulary we can get a single distribution over strings of any length
- **Why?**
  - because  $P(</s>|...)$  will be high enough that we are always guaranteed to stop after generating a finite number of words.

# LANGUAGE MODELING WITH N-GRAM

## Maximum likelihood estimate

$$P(w_i | w_{i-1}) = \frac{\text{count}(w_{i-1}, w_i)}{\text{count}(w_{i-1})}$$

$$P(w_i | w_{i-1}) = \frac{c(w_{i-1}, w_i)}{c(w_{i-1})}$$

## Example

$$P(w_i | w_{i-1}) = \frac{c(w_{i-1}, w_i)}{c(w_{i-1})}$$

<s> I am Sam </s>  
<s> Sam I am </s>  
<s> I do not like green eggs and ham </s>

$$P(\text{I} | \text{<s>}) = \frac{2}{3} = .67 \quad P(\text{Sam} | \text{<s>}) = \frac{1}{3} = .33 \quad P(\text{am} | \text{I}) = \frac{2}{3} = .67$$

## ESTIMATING PROBABILITIES - PRACTICAL ISSUES

- We do everything in log space
  - Avoid underflow
  - Computationally adding is faster than multiplying

$$\log(p_1 \times p_2 \times p_3 \times p_4) = \log p_1 + \log p_2 + \log p_3 + \log p_4$$

## NUMBER OF POSSIBLE PARAMETERS

- Estimating the number of parameter per n-gram language model.
- Given a vocabulary  $V$  of  $|V|$  unique tokens, where  $|V| = 10^4$ 
  - **Unigram** model:  $|V|$  parameters  $\Leftrightarrow 10^4$  parameters
  - **Bigram** model:  $|V|^2$  parameters  $\Leftrightarrow 10^8$  parameters
  - **Trigram** model:  $|V|^3$  parameters  $\Leftrightarrow 10^{12}$  parameters

# SHAKESPEARE AS CORPUS

- Number of words (symbols) = 884,647
- Tokens,  $V=29,066$
- Shakespeare produced 300,000 bigrams
- bigram types out of  $V^2= 844$  million possible bigrams
- So, 99.96% of the possible bigrams were never seen (*have zero entries in the table*)  
*844 million – 300,000 unused bigrams*
- Quadrigrams worse:
  - What's coming out looks like Shakespeare because it *is* Shakespeare

## EVALUATION: HOW GOOD IS OUR MODEL?

- Does our language model prefer good sentences to bad ones?
- Assign higher probability to “real” or “frequently observed” sentences
  - than “ungrammatical” or “rarely observed” sentences?
- We train parameters of our model on a **training set**.
- We test the model’s performance on data we haven’t seen.
  - A **test set** is an unseen dataset that is different from our training set, totally unused.
  - An **evaluation metric** tells us how well our model does on the test set.

# EXTRINSIC EVALUATION OF N-GRAM MODELS

- Best evaluation for comparing models A and B
- Embed each model in a task
  - spelling corrector,
  - speech recognizer,
  - Machine Translation system
- Run the task, get an accuracy for A and for B
  - How many misspelled words corrected properly
  - How many words translated correctly
- Compare accuracy for A and B

*Time-consuming; can take days,  
weeks or even months*

# INTRINSIC EVALUATION OF N-GRAM MODELS

- Sometimes use **intrinsic** evaluation: **perplexity**
- Bad approximation
  - unless the test data looks **just** like the training data
  - So **generally, only useful in pilot experiments**
- But is helpful to think about.



# PERPLEXITY

- Perplexity is the inverse probability of the test set, normalized by the number of words  $N$ :

$$\text{perplexity}(W) = P(x_1 x_2, \dots, x_N)^{-1/N} = \sqrt[N]{\frac{1}{P(x_1 x_2, \dots, x_N)}} = \sqrt[N]{\prod_{i=1}^N \frac{1}{P(x_i | x_1, \dots, x_{i-1})}}$$

- For unigram  $\text{perplexity}(W) = \sqrt[N]{\prod_{i=1}^N \frac{1}{P(x_i)}}$
- For bigram  $\text{perplexity}(W) = \sqrt[N]{\prod_{i=1}^N \frac{1}{P(x_i | x_{i-1})}}$

NB: Minimizing perplexity is the same as maximizing probability

# INTUITION OF PERPLEXITY


## ■ The Shannon Game:

- How well can we predict the next word?

I always order pizza with cheese and \_\_\_\_

The 33<sup>rd</sup> President of the US was \_\_\_\_

I saw a \_\_\_\_



laptop	0.09
mushrooms	0.02
pepperoni	0.1
anchovies	0.0001
fried rice	0.005
and	1e-10
cook	2e-15

- Unigrams are terrible at this game. (Why?)
- ## ■ A better language model is the one that assigns a higher probability to the most appropriate word

## LIMITATIONS – STORAGE PROBLEMS

### Storage Problem:

The need to store count for all n-grams you saw in the corpus.

$$P(\mathbf{w} \mid I \text{ want to eat}) =$$

$$\frac{\text{count}(I \text{ want to eat } \mathbf{w})}{\text{count}(I \text{ want to eat})}$$

Increasing  $n$  or increasing corpus size  $\Leftrightarrow$  increases model size

# LIMITATIONS - SPARSITY PROBLEM

## Sparsity Problem 1:

- What if “*I want to w*” never occurred in the corpus?
- Then *w* has probability 0

**Solution - smoothing:** Add small  $\delta$  to the count for every  $w \in V$ .

*count(waiting for the **w**)*

$P(\mathbf{w} | \text{waiting for the}) =$

*count(waiting for the)*

## Sparsity Problem 2:

- What if “*I want to*” never occurred in the corpus?
- Then we cannot calculate probability of any *w*.

**Solution - backoff:**  
condition on the (n-1)-gram,  
i.e. “*I want*”

Larger  $n$  makes sparsity problem worse. Typically  $n$  should be less than or equal to 5

# THE PERILS OF OVERFITTING - SPARSITY

N-grams only work well for word prediction if the test corpus looks like the training corpus.

- In real life, it often doesn't
  - We need to train robust models that generalize!
- One kind of generalization: Zeros!
  - Things that don't ever occur in the training set but occur in the test set

# SMOOTHING METHODS

## Smoothing methods

- Additive smoothing
- Good-Turing estimate
- Jelinek-Mercer smoothing (interpolation)
- Katz smoothing (backoff)
- Witten-Bell smoothing
- Absolute discounting
- Kneser-Ney smoothing

# ADDITIVE SMOOTHING

- Idea: pretend we've seen each n-gram  $\delta$  times more than we have.
- Typically,  $0 < \delta \leq 1$ .
- Lidstone and Jeffreys advocate  $\delta = 1$ .
- Gale & Church (1994) argue that this method performs poorly.

$$p_{add}(w_i | w_{i-n+1}^{i-1}) = \frac{\delta + c(w_{i-n+1}^i)}{\delta|V| + \sum_{w_i} c(w_{i-n+1}^i)}$$

# ADD-ONE ESTIMATION

- Also called Laplace smoothing
- Pretend we saw each word one more time than we did
- Just add one to all the counts!
- MLE estimate:

$$P_{MLE}(w_i | w_{i-1}) = \frac{c(w_{i-1}, w_i)}{c(w_{i-1})}$$

- Add-1 estimate:

$$P_{Add-1}(w_i | w_{i-1}) = \frac{c(w_{i-1}, w_i) + 1}{c(w_{i-1}) + V}$$



# BACKOFF AND INTERPOLATION

- Sometimes it helps to use **less** context
  - Condition on less context for contexts you haven't learned much about
- **Backoff:**
  - use trigram if you have good evidence,
  - otherwise, bigram, otherwise unigram
- **Interpolation:**
  - mix unigram, bigram, trigram
- NB: Interpolation works better

# REFERENCES